



SAPIENZA
UNIVERSITÀ DI ROMA

Data Mining Technology for Business and Society

Homework 1

Tran Luong Bang - *tran.1956419@studenti.uniroma1.it*

Juan Mata Naranjo - *matanaranjo.1939671@studenti.uniroma1.it*

1 Part 1.1

Before we deep dive into each of the individual datasets and their respective results we will start out by first presenting the number of documents we have, the number of queries applied over these documents and finally the number of queries for which we have ground truth information at our disposal. When computing the evaluation metrics we will assume that the queries for which we don't have ground truth don't exist (we remove them since we cannot compare them against anything)¹.

	Cranfield Dataset	Time Dataset
Num Indexed Docs	1400	423
Num Queries	225	83
Num Queries in GT	110	80

Table 1: Overview Table

1.1 Cranfield Dataset:

The configurations we have tested are the following:

Conf. ID	Text Analyzer	Scoring Functions
1	RegexTokenizer() LowercaseFilter() StopFilter(stoplist = STOP_WORDS)	scoring.MultiWeighting(scoring.Frequency(), title=scoring.BM25F(), content=scoring.TF_IDF())
2	RegexTokenizer() StopFilter(stoplist = STOP_WORDS) LowercaseFilter() StemFilter()	scoring.TF_IDF()
3	StemmingAnalyzer(stoplist=STOP_WORDS)	scoring.TF_IDF()
4	FancyAnalyzer()	scoring.BM25F(K1=2, B=.8)
5	SimpleAnalyzer()	scoring.TF_IDF()
6	StandardAnalyzer()	scoring.Frequency()
7	FancyAnalyzer()	scoring.MultiWeighting(scoring.Frequency(), title=scoring.TF_IDF(), content=scoring.BM25F())
8	StemmingAnalyzer(stoplist=STOP_WORDS)	scoring.BM25F()
9	RegexTokenizer() StopFilter(stoplist = STOP_WORDS) LowercaseFilter() StemFilter()	scoring.MultiWeighting(scoring.Frequency(), title=scoring.TF_IDF(), content=scoring.BM25F(K1=2, B=.8))
10	SimpleAnalyzer()	scoring.BM25F(K1=1.2, B=.75)
11	StemmingAnalyzer(stoplist=STOP_WORDS)	scoring.TF_IDF()
12	FancyAnalyzer()	scoring.TF_IDF()

Table 2: Configuration Overview for Cranfield Dataset

For which the resulting MRR and R-Precision results are the following (the table presented below is already sorted from highest to lowest MRR values):

¹All code can also be found under the following git repository: https://github.com/JuanMataNaranjo/DMT_HW1

Conf. ID	MRR	Mean	Min	1st Quartile	Median	3rd Quartile	Max
8	0.458	0.237	0.000	0.000	0.186	0.389	1.000
9	0.453	0.230	0.000	0.000	0.167	0.379	1.000
10	0.417	0.193	0.000	0.000	0.143	0.282	1.000
4	0.415	0.216	0.000	0.000	0.186	0.333	1.000
7	0.412	0.212	0.000	0.000	0.170	0.333	1.000
2	0.368	0.165	0.000	0.000	0.106	0.250	1.000
3	0.368	0.165	0.000	0.000	0.106	0.250	1.000
11	0.368	0.165	0.000	0.000	0.106	0.250	1.000
1	0.323	0.147	0.000	0.000	0.025	0.250	1.000
12	0.311	0.137	0.000	0.000	0.000	0.222	1.000
6	0.248	0.104	0.000	0.000	0.000	0.167	0.833
5	0.126	0.056	0.000	0.000	0.000	0.000	0.667

Table 3: MRR and R-Precision Overview for Cranfield

Finally, the next figure shows the results for the NCDG@k and P@k metrics for $k \in [1, 3, 5, 10]$

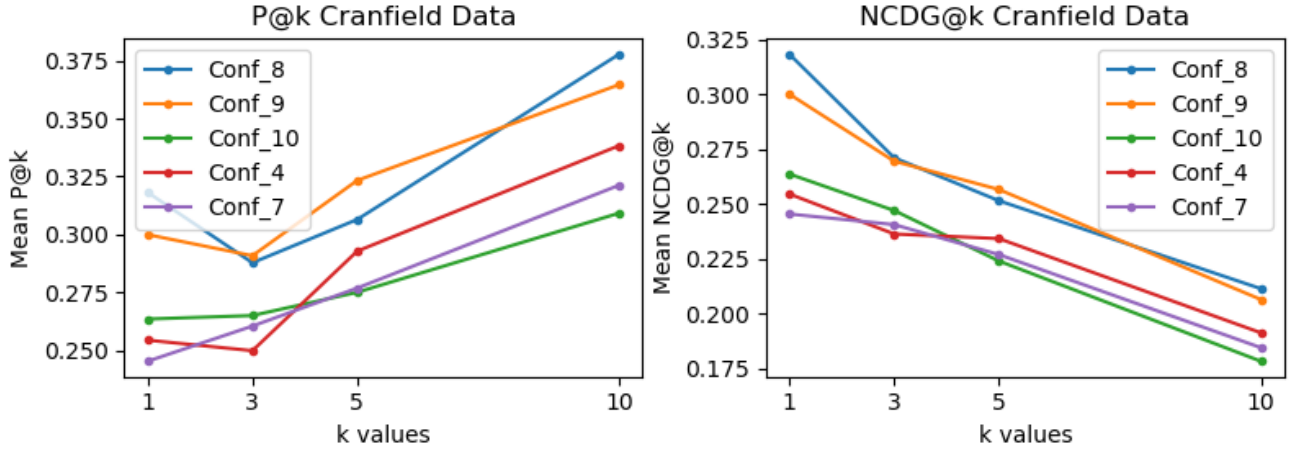


Figure 1: P@k and NCDG@k Plot Results for Cranfield Dataset

If we were to choose the best Search Engine solely based on the NCDG@k curve we would choose the configuration for which this ratio is highest when k is maximal. Therefore, based on the previous plot, the champion configuration would be **Conf 8** (Stemming Analyzer & BM25F).

1.2 Time Dataset:

We will follow the same structure and order to present the Time Dataset results:

Conf. ID	Text Analyzer	Scoring Functions
1	RegexTokenizer() LowercaseFilter() StopFilter(stoplist = STOP_WORDS)	scoring.TF.IDF()
2	RegexTokenizer() StopFilter(stoplist = STOP_WORDS) LowercaseFilter() StemFilter()	scoring.BM25F(K1=1.2, B=.75)
3	StemmingAnalyzer(stoplist=STOP_WORDS)	scoring.TF.IDF()
4	FancyAnalyzer()	scoring.BM25F(K1=2, B=.8)
5	SimpleAnalyzer()	scoring.TF.IDF()
6	StandardAnalyzer()	scoring.Frequency()
7	FancyAnalyzer()	scoring.BM25F()
8	StemmingAnalyzer(stoplist=STOP_WORDS)	scoring.BM25F()
9	RegexTokenizer() StopFilter(stoplist = STOP_WORDS) LowercaseFilter() StemFilter()	scoring.BM25F(K1=2, B=.8)
10	SimpleAnalyzer()	scoring.BM25F(K1=1.2, B=.75)
11	StemmingAnalyzer(stoplist=STOP_WORDS)	scoring.BM25F(K1=2, B=.8)
12	FancyAnalyzer()	scoring.TF.IDF()

Table 4: Configuration Overview for Time Dataset

Conf. ID	MRR	Mean	Min	1st Quartile	Median	3rd Quartile	Max
11	0.725	0.555	0.000	0.264	0.539	0.906	1.000
8	0.717	0.562	0.000	0.200	0.564	1.000	1.000
2	0.708	0.557	0.000	0.200	0.564	1.000	1.000
9	0.703	0.550	0.000	0.200	0.539	0.906	1.000
4	0.689	0.539	0.000	0.200	0.500	0.889	1.000
7	0.678	0.549	0.000	0.321	0.500	0.889	1.000
10	0.669	0.542	0.000	0.321	0.500	0.889	1.000
1	0.590	0.351	0.000	0.000	0.376	0.600	1.000
3	0.566	0.309	0.000	0.000	0.333	0.500	1.000
12	0.536	0.285	0.000	0.000	0.171	0.510	1.000
6	0.457	0.238	0.000	0.000	0.063	0.500	1.000
5	0.240	0.084	0.000	0.000	0.000	0.125	1.000

Table 5: MRR and R-Precision Overview for Time

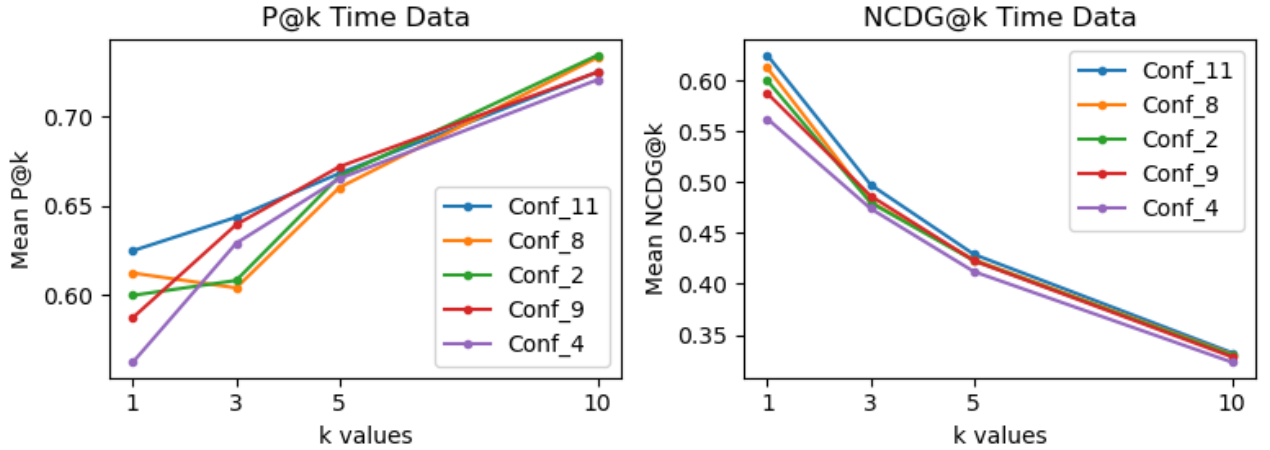


Figure 2: P@k and NCDG@k Plot Results for Time Dataset

2 Part 1.2

To select the best Search-Engine modules for the “AwesomeSocialApp”, we will use the following evaluation metrics for ranking problems:

Precision@k: Considering the definition, which states *of all the query results we predicted to be True, how many were actually True?*. The graph shows that SE1 performs slightly better than SE3 and much better than SE2

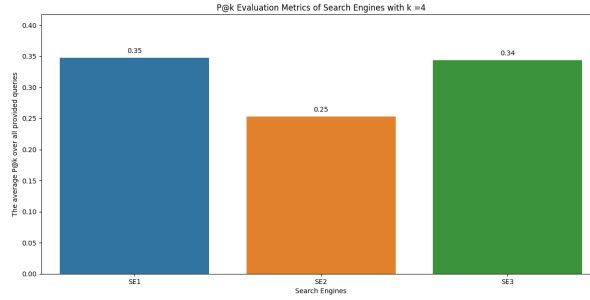


Figure 3: P@k Evaluation Metrics of Search Engines

Recall@k: Similarly *of all the query results that were actually True, how many were predicted to be True*. As we can see on graph below, SE3 should be selected compared with SE1 and SE2

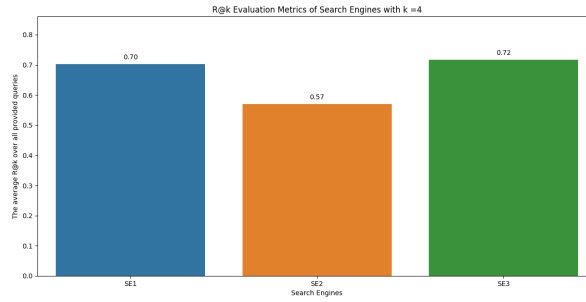


Figure 4: R@k Evaluation Metrics of Search Engines

nCDG@k: Normalized Discounted Cumulative Gain is a metric that is particularly good at evaluating search engines which return ranked outputs (the order of the output is relevant for this metric)

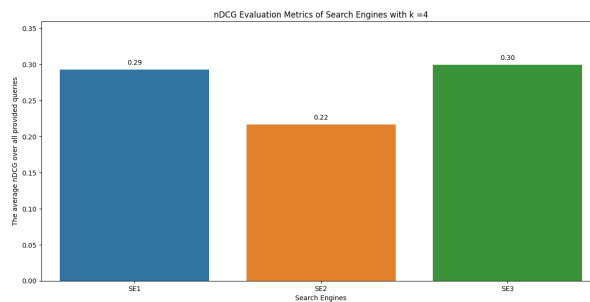


Figure 5: nDCG@k Evaluation Metrics of Search Engines

As we can see the figures on 3 graphs, SE1 and SE3 have a similar performance and much better than SE 2. But SE3 have a better performance on **nDCG@k** and **Recall@k** evaluation metrics, we highly recommend The “DummyDataScience” company to choose *SE3* for this application.

3 Part 2.1

The first thing we have to do before running the LSH algorithm is to figure out the best r (number of rows in each band), b (number of bands) and n (number of hash functions used) parameters such that we can minimize the number of False Positives and False Negatives. In particular, the choice of parameters will mainly help us to minimize the quantity of False Negatives. For the parameter choices we also need to take into account the following constraints:

$$r \cdot b = n \quad (1)$$

$$0.97 > 1 - (1 - 0.95^r)^b \quad (2)$$

From the second constraint, solving for b we can see that the following relation holds:

$$b > \frac{\log(1 - 0.97)}{\log(1 - 0.95^r)} \quad (3)$$

Therefore, to minimize the False Negatives we will stress the 0.97 bound as much as possible, specifically up to 0.999. This will reduce the number of False Negatives, but will of course increase the number of False Positives. Since we will be able to remove some of the False Positives later on we are willing to make this compromise. So finally, fixing a value of $r = 15$ we have decided to use the following parameters:

r	15
b	12
n	180

Table 6: Parameter Overview

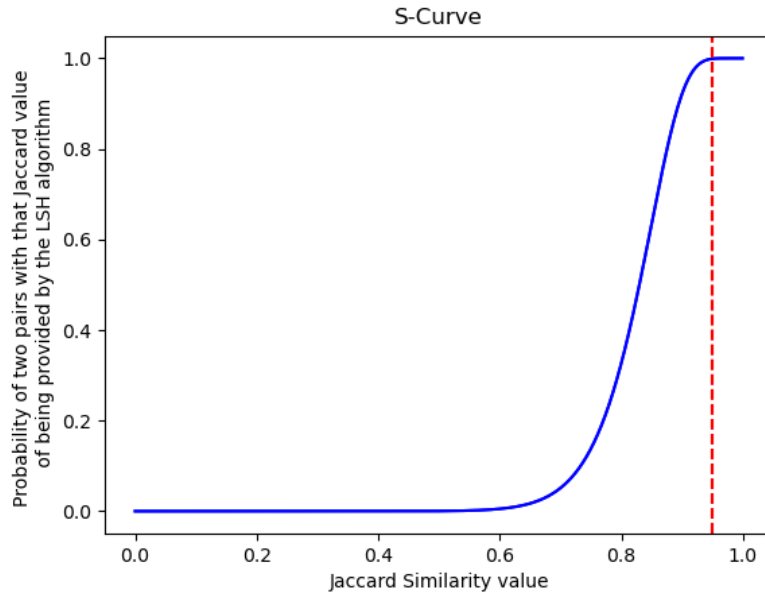


Figure 6: Relationship between Jaccard Similarity and probability that LSH will return that pair as a potential near duplicate

After running the LSH algorithm with the previous set of parameters we got the following number of *potential* near duplicate matches: **34655**. However, as we already mentioned before, we can still reduce the number of False Positives a posteriori by computing the Jaccard Similarity over the set of pairs provided by the LSH tool, and therefore eliminating those pairs which we will not be consider as near duplicates. For this purpose we have used the

code named "False Positive Reduction" (in which we re-compute the real Jaccard similarity and not the approximated one), after which the number of *final* near duplicates is: **33868**.

The total time required to run the LSH tool with the parameters highlighted previously was of **7 minutes and 15 seconds** on our machine.

4 Part 2.2

We are looking for all near-duplicate title of song in the dataset with 250K song, it means that there are more than 3.12×10^{20} pairs of song titles. Therefore to solve this problem, we have applied the Near duplicate Detection with the purpose of reducing the number of pairs of song titles to check.

- Before shingling, we have preprocessed the data (e.g. replacing "-" with " ")
- The Jaccard similarity between their associated sets of shingles is **1**
- Number of hash fuctions is $n = 200$, with number of bands is $b = 20$ and number of rows $r = 10$ per band.

With these assumptions, by applying the NDD, we have found **435515** pairs of near-duplicated song titles. It is also relevant to highlight that in this case, the Jaccard Similarity is exactly equal to 1, so we can not reduce the number of False Negative by changing the parameters like r and b.

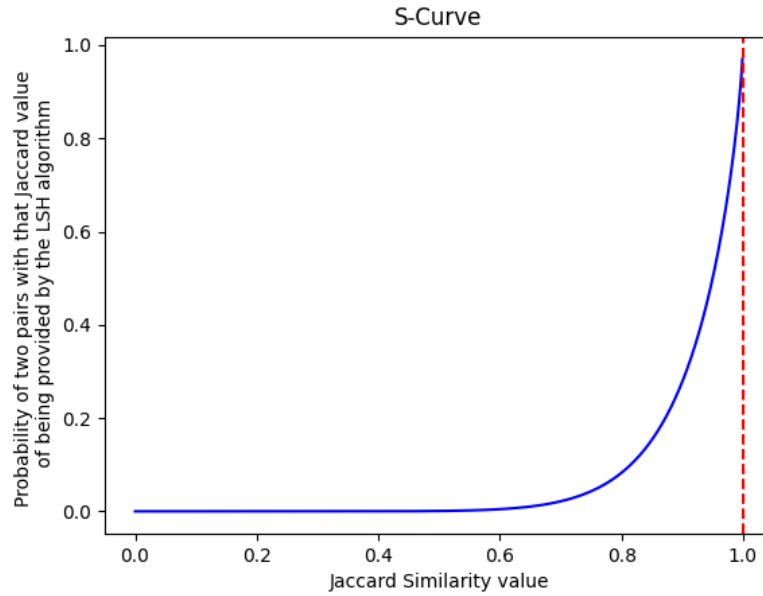


Figure 7: Relationship between Jaccard Similarity and probability that LSH will return that pair as a potential near duplicate. We see that when Jacc=1 the LSH output is surely also 1

To reduce the number of False Positive, we have tried to apply 2 methods:

1. Reducing the Probability of 2 pairs with that Jaccard value provided by LSH algorithm. We have got the same number of pairs **435515**
2. Computing the Jaccard Similarity on the sets of pairs generated by applying LSH tool, and therefore we can remove those pairs not considered as near-duplicate. After excuting it, we found the same number of pairs **435515**

It means that there is no Fasle Positive found and all pairs of song titles we found are near-duplicate.

In conclusion, after applying this method, we have found **435515 pairs of near-duplicate song titles** with around **2 hours** executing on our computer.